

# Pylog: Prolog in Python

Russ Abbott, Jungsoo Lim  
Department of Computer Science  
California State University, Los Angeles  
Los Angeles, California 90032

Email: rabbott@calstatela.edu, jlim34@calstatela.edu

Jay Patel  
Visa  
Foster City  
USA

Email: imjaypatel12@gmail.com

**Abstract—Context:** What is the broad context of the work?

What is the importance of the general research area?

Pylog inhabits three programming contexts.

- Pylog explores the integration of two distinct programming language paradigms: (i) the modern general purpose programming paradigm, including features of procedural programming, object-oriented programming, functional programming, and meta-programming, here represented by Python, and (ii) logic programming, whose primary features are logic variables (and unification) and built-in depth-first backtracking search, here represented by Prolog. These logic programming features are generally missing from modern general purpose languages. Pylog illustrates how these two features can be implemented in and integrated into Python.
- Pylog demonstrates the breadth and broad applicability of Python. Although Python is one of the most widely used programming languages for teaching introductory programming, it has also become very widely used for sophisticated programming tasks. One of the reasons for its popularity is the range of capabilities it offers—most of which are not used in elementary programming classes. Pylog makes effective use of many of those capabilities.
- Pylog exemplifies programming at its best. Pylog is first-of-all a programming exercise: How can the primary features of logic programming be integrated with Python? Secondly, Pylog uses features of Python in ways that are both intended and innovative. These include distinguishing between two uses of Python's for-loop structure—as choicepoints and as aggregating constructs. The overall result is software worth reading.

**Inquiry:** What problem or question does the paper address? How has this problem or question been addressed by others (if at all)?

The primary issue addressed is how logic variables and backtracking can be integrated cleanly into a Python framework. Although significant work has been done in this area, much of it well done, most has been incomplete. Pylog is the first complete system (as far as we know) to achieve the goal of full integration. Also, as far as we know, this paper offers the first thorough explanation for how such integration can be accomplished.

**Approach:** What was done that unveiled new knowledge?

Pylog demonstrates how logic variables and backtracking can be interwoven with standard Python data structures and control structures.

**Knowledge:** What new facts were uncovered? If the research was not results oriented, what new capabilities are enabled by the work?

Pylog is available as a library for use in Python software. Pylog's implementation techniques and insights may be used in Python programs not limited to logic programming.

**Grounding:** What argument, feasibility proof, artifacts, or results

and evaluation support this work?

By its existence Pylog demonstrates that logic variables and backtracking can be integrated into Python.

**Importance:** Why does this work matter?

Python is known to be compatible with functional programming and other paradigms. This work shows that it is also compatible with logic programming. This work demonstrates the power and elegance of well-designed software.

The Pylog code is available at [this GitHub repository](#).

ACM CCS 2012

- Software and its engineering ~ Multiparadigm languages;

## I. INTRODUCTION

The birth announcement for Artificial Intelligence took the form of a 1955 workshop proposal. The proposal claimed that every aspect of learning—or any other feature of intelligence—can in principle be so precisely described that a machine can be made to simulate it.[15]

Allen Newell's and Herbert Simon's demonstration of their Logic Theorist seemed to fulfill that promise. Newell and Simon claimed that their program *solved the venerable mind-body problem*. [21] A year later Simon doubled-down.

[T]here are now in the world machines that can think, that can learn, and that can create. Moreover, their ability to do these things is going to increase rapidly until—in a visible future—the range of problems they can handle will be coextensive with the range to which the human mind has been applied. [25]

Perhaps not unexpectedly, optimism about the power of symbolic AI, as this work was (as is still) known faded into the gloom of what has been labelled the AI winter.

But winter was followed by spring, which brought the sprouting of work on (non-symbolic) deep neural networks. Andrew Ng said of that development,

Just as electricity transformed almost everything 100 years ago, today I actually have a hard time thinking of an industry that I don't think AI will transform in the next several years. [18]

But another disappointment followed. Deep neural nets are surprisingly susceptible to what are known as *adversarial* attacks. Small perturbations to images

that are (almost) imperceptible to human vision can cause a neural network to completely change its prediction—e.g., a minimally modified correctly classified image of a school bus is classified as an ostrich. Even worse, the classifiers report high confidence on the wrong prediction.[1]

(Adversarial images did not kill off deep learning. They have been co-opted, and their use is now built into deep neural network training methodologies.[24])

In the meantime, work on symbolic AI was not abandoned. Over the past few decades, symbolic AI has evolved into what is now known as constraint programming.

Constraint programming was born in the 1980 as an outgrowth of the sudden interest in logic programming triggered by the Japanese fifth generation initiative.[23] Logic programming led to Constraint Logic Programming which evolved into constraint programming. (A familiar constraint programming problem is the well-known  $n$ -queens problem: how can you place  $n$  queens on an  $n \times n$  chess board so that no queen threatens any other queen? There are, of course, many practical applications of constraint programming as well.)

In 1997, Eugene Frueder characterized constraint programming as *the Holy Grail of computer science: the user simply states the problem and the computer solves it*. [11]

Software that solves constraint programming problems are known as solvers. Significant progress has been made on solver development over the past few decades.

As a sub-discipline of computer science constraint programming has a number of nice properties. Its structure and limits are well understood: there will be no grand disappointments similar to those that followed the birth of artificial intelligence—unless quantum computing, once implemented, turns out to be a bust. Constraint programming is closely related to computational complexity, which provides a well-studied theoretical framework. Nor will there be negative surprises such as adversarial images that rock the discipline.

Furthermore, the fundamentals of solver technology are well-understood. Solving a constraint problem is a search exercise: find values for the uninstantiated variables that satisfy the constraints.

Improvements are generally incremental and consist primarily of improved search heuristics. For example, in the  $n$ -queens problem one can propagate solution steps by marking as unavailable board squares that are attacked by each newly placed piece. Propagation reduces search times significantly. We will see simple example heuristics below.

Constraint programming solvers are now available in multiple forms. The language [MiniZinc](#) allows users to express constraints in what is essentially executable predicate calculus.

Solvers are also available as package add-ons to many programming languages: [Choco](#) and [JaCoP](#) (two Java libraries), [Iz-C](#) (a C library), [Oscar/CBLS](#) and [Yuck](#) (two Scala libraries), and Google's [OR-tools](#) (a collection of C++ libraries, which sport Python, Java, and .NET front ends).

In both approaches, the solver is relatively isolated from the rest of the programming environment. One sets up a problem,

either directly in predicate calculus (as in [MiniZinc](#)) or in the host language, and then calls on the solver to solve it. There is little back-and-forth between code in the host language and the solver package. This is fine for one-shot problems, but some problems call for more interaction with the solver.

There is significantly more integration when working in a language like [Picat](#), a languages that combines features of logic programming and imperative programming, or when starting with, for example, [SICStus Prolog](#) or [SWI Prolog](#) and adding a Finite Domain package. Unfortunately, a logic programming background is a prerequisite to working with such systems.

A commonly discussed—although less frequently implemented—alternative is known as *shallow embedding*[13], [12]. One implements solver capabilities directly in a host language in such a way that programs in the host-language can use the solver capabilities as needed. Two recent examples are [Kanren](#)[20], logic programming in Python, and [Muli](#)[8], an embedding of constraint-logic capabilities in Java.

Although working with shallowly embedded solvers sounds attractive, neither [Kanren](#) nor [Muli](#) are idea candidates. [Kanren](#) is open source, but no implementation documentation is available. [Muli](#) is much more forthcoming. [Dageförde](#) and [Kuchen](#) describe the [Muli](#) virtual machine[7]. But the documentation is fairly technical. Many of those who would like to use [Muli](#) may have a difficult time working their way through it.

Our goal in this paper is to offer an under-the-covers tutorial to those who may want to write software that works interactively with a shallowly embedded solver—or to those who perhaps may even want to build their own.

We describe [Pylog](#), a shallow embedding of the core Prolog capabilities in Python. Our focus will be on helping the reader understand how Prolog's two fundamental features, backtracking and logic variables, can be implemented in standard Python—as well as in other languages with similar capabilities. [Pylog](#) is implemented in Python and should be accessible to anyone reasonable fluent in Python.

## II. RELATED WORK

Quite a bit of work has been done in implementing Prolog features in Python. As far as we can tell, none of it is as complete and as fully thought through as [Pylog](#). But nearly all make important contributions. Following, in chronological order, are the authors' own descriptions—lightly edited for clarity and brevity.

- [Berger](#) (2004) [3]. [Pythologic](#).

Python's meta-programming features are used to enable the writing of functions that include Prolog-like features.

- [Bolz](#) (2007) [4] [A Prolog Interpreter in Python](#).

A proof-of-concept implementation of a Prolog interpreter in [RPython](#), a restricted subset of the Python language intended for system programming. Performance compares reasonably well with other embedded Prologs.

- [Delford](#) (2009) [9] [Pylog](#).

A proof-of-concept implementation of a Prolog interpreter in RPython.

- Frederiksen (2011) [10] Pike.  
A form of Logic Programming that integrates with Python.
- Meyers (2015) [16] Prolog in Python.  
A hobby project developed over a number of years.
- Maxime (2016) [14] Prology: Logic programming for Python3.  
A minimal library that brings Logic Programming to Python.
- Piumarta (2017) [19] Notes and slides from a course on programming paradigms. Pylog started as a fork of Piumarta's work. Unfortunately it is no longer available.
- Thompson (2017) [26] Yield Prolog.  
Enables the embedding of Prolog-style predicates directly in Python.
- Santini (2018) [22] The pattern matching for python you always dreamed of.  
Pampy is small, reasonably fast, and often makes code more readable.
- Cesar (2019) [6] Prol: a minimal Prolog interpreter in a few lines of Python.
- Miljkovic (2019)[17] A simple Prolog Interpreter in a few lines of Python 3.
- Rocklin (2019) [20] kanren: Logic Programming in Python.  
Enables the expression of relations and the search for values that satisfy them. A Python implementation of miniKanren [5]

Much of the preceding work is quite well done. As this survey suggests, most of the important ideas for embedding Prolog-like capabilities in Python have been known for a while. Pylog's goal is to be a more fully developed, more fully explained, and more integrated version of these ideas.

### III. FROM PYTHON TO PROLOG (LISTINGS IN APPENDIX ??)

This section offers a reasonably detailed overview of Pylog and how it relates to Prolog. Our strategy is to show how a standard Python program can be transformed, step-by-step, into a structurally similar Prolog program.

As an example problem, we use the computation of a transversal. Given a sequence of sets (in our case lists without repetition), a transversal is a non-repeating sequence of elements with the property that the  $n^{th}$  element of the traversal belongs to the  $n^{th}$  set in the sequence. For example, the sets<sup>1</sup>  $[[1, 2, 3], [2, 4], [1]]$  has three transversals:  $[2, 4, 1]$ ,  $[3, 2, 1]$ , and  $[3, 4, 1]$ . We use the transversal problem because it

<sup>1</sup>From here on, we refer informally to the lists in our example as *sets*.

lends itself to depth-first search, the default Prolog control structure.<sup>2</sup>

We will discuss five functions for finding transversals—the first four in Python, the final one in standard Prolog. As we discuss these programs we will introduce various Pylog features. Here is a road-map for the programs to be discussed and the Pylog features they illustrate. (To simplify formatting, we use *tvsl* in place of *transversal*)

- 1) *tvsl\_dfs\_first* is a standard Python program that performs a depth-first search. It returns the first transversal it finds. It contains no Pylog features, but it illustrates the overall structure the others follow.
- 2) *tvsl\_dfs\_all*. In contrast to *tvsl\_dfs\_first*, *tvsl\_dfs\_all* finds and returns *all* transversals. A very common strategy, and the one *tvsl\_dfs\_all* uses, is to gather all transversals into a collection as they are found and return that collection at the end.
- 3) *tvsl\_yield* also finds and returns all transversals, but it returns them one at a time as requested, as in Prolog. *tvsl\_yield* does this through the use of the Python generator structure, i.e., the **yield** statement. This moves us an important step toward a Prolog-like control structure.
- 4) *tvsl\_yield\_lv* introduces logic variables.
- 5) *tvsl\_prolog* is a straight Prolog program. It is operationally identical to *tvsl\_yield\_lv*, but of course syntactically very different.

The first three Python programs have similar signatures.

<pre>def tvsl_python_1_2_3(sets: List[List[int]],</pre>	<pre>partial_tr : Tuple ) -&gt; &lt; some return type &gt;:</pre>
---	---

(The return types differ from one program to an other.)

Both the fourth Python program and the Prolog program have a third parameter. Their return type, if any, is not meaningful. In these programs, transversals, when found, are returned through the third parameter—as one does in Prolog.

The signatures all have the following in common.

- 1) The first argument lists the sets for which a transversal is desired, initially the full list of sets. The programs recursively steps through the list, selecting an element from each set. At each recursive call, the first argument lists the remaining sets.

<sup>2</sup>We use traditional, i.e., naive, depth-first search. Most modern Prologs include a constraint processing package such as CLP(FD)[27], which makes search much more efficient.

Application of such constraint rules eliminates much of the backtracking inherent in naive depth-first search. Powerful as they are, we do not use such techniques in this paper.

- 2) The second argument is a partial transversal consisting of elements selected from sets that have already been scanned. Initially, this argument is the empty tuple.<sup>3</sup>
- 3) The third parameter, if there is one, is the returned transversal.
  - a) The first three programs have no third parameter. They return their results via **return** or **yield**.
  - b) The final Python function and the Prolog predicate both have a third parameter. Neither uses **return** or **yield** to return values. In both, the third argument, initially an uninstantiated logic variable, is unified with a transversal if found.

We now turn to the details of the programs. For each program, we introduce the relevant Python/Pylog constructs and then discuss how they are used in that program.

#### A. *tvsl\_dfs\_first* (Listings in Appendix ??)

*tvsl\_dfs\_first* uses standard depth-first search to find a single transversal. As Listing ?? shows, when we reach the end of the list of sets (line 3), we are done. At that point we return *partial\_transversal*, which is then known to be a complete transversal.

The return type is *Optional[Tuple]*,<sup>4</sup> i.e., either a tuple of *ints*, or **None** if no transversal is found. The latter situation occurs when, after considering all elements of the current set (*sets[0]*) (line 6), we have not found a complete transversal.

It may be instructive to run *transversal\_dfs\_first*

```
sets = [[1, 2, 3], [2, 4], [1]]
transversal_dfs_first(sets)
```

and look at the log (Listing ??) created by the *Trace* decorator.<sup>5</sup>

The log (Listing ??) shows the value of the parameters at the start of each function execution. When *sets* is the empty list (line 3), we have found a transversal—which the *Trace* function indicates with `<=`. On the other hand, when the function reaches a dead-end, it “backtracks” to the next element in the current set and tries again.

The first three lines of the log show that we have selected (1, 2) as the *partial\_transversal* and must now select an element of [1], the remaining set. Since 1 is already in the *partial\_transversal*, it can’t be selected to represent the final set. So we (blindly, as is the case with naive depth-first search) backtrack (line 6 in the code) to the selection from the second set. We had initially selected 2. Line 4 of the log shows that we have now selected 4. Of course that doesn’t help.

Having exhausted all elements of the second set, we backtrack all the way to our selection from the first set (again line 6 in the code). Line 5 of the log shows that we have now selected 2 from the first set and are about to make a selection from the second set. We cannot select 2 from the second set

since it is already in the *partial\_transversal*. Instead, we select 4 from the second set. We are then able to select 1 from the final set, which, as shown on line 7, completes the transversal.

Even though this is a simple depth-first search, it incorporates (what appears to be) backtracking. What implements the backtracking? In fact, there is no (explicit) backtracking. The nested **for**-loops produce a backtracking effect. Prolog, uses the term *choicepoint* for places in the program at which (a) multiple choices are possible and (b) one wants to try them all, if necessary. Pylog implements choicepoints by means of such nested **for**-loops and related mechanisms.

#### B. *for*-loops as choice points and as computational aggregators (Listings in Appendix ??)

Although we are using a standard Python **for**-loop, it’s worth noticing that in the context of depth-first search, a **for**-loop does, in fact, implement a choicepoint. A choicepoint is a place in the program at which one selects one of a number of options and then moves forward with that selection. If the program reaches a dead-end, it “backtracks” to the choicepoint and selects another option. That’s exactly what the **for**-loop on line 6 of the program does: it generates options until either we find one for which the remainder of the program succeeds, or, if the options available at that choicepoint are exhausted, the program backtracks to an earlier choicepoint.

Is there a difference between this way of using **for**-loops and other ways of using them? One difference is that with traditional **for**-loops, e.g., one that would be used in a program to find, say, the largest element of a list (without using a built-in or library function like *max* or *reduce*), each time the **for**-loop body executes, it does so in a context produced by previous executions of the **for**-loop body.

Consider the simple program *find\_largest*, (Listing ??), which finds the largest element of a list. The value of *largest* may differ from one execution of the **for**-loop body to the next. No similar variables appear in the **for**-loop body of *tvsl\_dfs\_first*.

The **for**-loop in *find\_largest* performs what one might call computational aggregation—results aggregate from one execution of the **for**-loop body to the next. In contrast, the **for**-loop in *tvsl\_dfs\_first* leaves no traces; there is no aggregation from one execution of the body to the next.

In addition, **for**-loops that function as a choicepoint define a context within which the selection made by the **for**-loop holds. Of course, the variables set by any **for**-loop are generally limited to the body of the **for**-loop. But **for**-loops that serve as choicepoints function more explicitly as contexts. Even when a choicepoint-type **for**-loop has only one option, it limits the scope of that option to the **for**-loop body. We will see examples in Section V-B Listings ?? and ?? when we discuss the *unify* function.

Most of the **for**-loops in this paper function as choicepoints. In particular, the **for**-loops in *tvsl\_dfs\_first*, *tvsl\_yield*, and *tvsl\_yield\_lv* all function as choicepoints. The **for**-loop in *tvsl\_dfs\_all* functions as an aggregator, aggregating transversals in the *all\_transversals* variable.

<sup>3</sup>The second parameter is of type *tuple* so that we can define an empty tuple as a default.

<sup>4</sup>There seems to be no way to specify a Tuple of arbitrary length, with *int* elements

<sup>5</sup>Code for the *Trace* decorator is included in Section ?? in the Appendix.



### C. *tvsl\_dfs\_all* (Listings in Appendix ??)

*tvsl\_dfs\_all* (Listing ??) finds and returns *all* transversals. It has the same structure as *tvsl\_dfs\_first* except that instead of returning a single transversal, transversals are added to *all\_transversals* (line 9), which is returned when the program terminates.

The following code segment produces the expected output. (Listing ?? shows a log.)

```
sets = [[1, 2, 3], [2, 4], [1]]
all_transversals = tvsl_dfs_all(sets)
print( \n All transversals: ', all_transversals)
```

If no transversals are found, *tvsl\_dfs\_all* returns an empty list.

### D. *tvsl\_yield* (Listings in Appendix ??)

*tvsl\_yield* (Listing ??), although quite similar to *tvsl\_dfs\_first*, takes a significant step toward mimicking Prolog. Whereas *tvsl\_dfs\_first* **returns** the first transversal it finds, *tvsl\_yield* **yields** *all* the transversals it finds—but one at a time.

Instead of looking for a single transversal as on lines 8 - 10 of *tvsl\_dfs\_first* and then **returning** those that are not **None**, *tvsl\_yield* uses **yield from** (line 8) to search for and **yield all** transversals—but only on request.

With *tvsl\_yield* one can ask for all transversals as follows.

```
sets = [[1, 2, 3], [2, 4], [1]]
for Transversal in tvsl_yield(sets):
    print(f Transversal: {Transversal} )
```

A full trace is shown in Listing ?? . This is discussed in more detail in Section IV.

### E. *tvsl\_yield\_lv* (Listings in Appendix ??)

*tvsl\_yield\_lv* (Listing ??) moves toward Prolog along a second dimension—the use of logic variables.

One of Prolog's defining features is its logic variables. A logic variable is similar to a variable in mathematics. It may or may not have a value, but once it gets a value, its value never changes—i.e., logic variables are immutable.

The primary operation on logic variables is known as *unification*. When a logic variable is *unified* with what is known as a *ground term*, e.g., a number, a string, etc., it acquires that term as its value. For example, if *X* is a logic variable,<sup>6</sup> then after *unify*(3, *X*), *X* has the value 3.

One can run *tvsl\_yield\_lv* as follows.

```
# Since we are using Pylog's logic variables, the input
# must be in that form.
sets = [PyList([1, 2, 3]), PyList([2, 4]), PyList([1])]
# Var() creates an uninstantiated logic variable
Complete_Transversal = Var()
for _ in tvsl_yield_lv(sets, PyTuple()),
    Complete_Transversal:
    print(f Transversal: {Complete_Transversal}\n )
```

The output, Trace included, will be as shown in Listing ??.

<sup>6</sup>The Python convention is to use only lower case letters in identifiers other than class names. Prolog requires that the first letter of a logic variable be upper case. In *tvsl\_yield\_lv* we use upper case letters to begin identifiers that refer to logic variables.

A significant difference between *tvsl\_yield* and *tvsl\_yield\_lv* is that in the **for**-loop that runs *tvsl\_yield*, the result is found in the loop variable, *Transversal* in this case. In the **for**-loop that runs *tvsl\_yield\_lv*, the result is found in the third parameter of *tvsl\_yield\_lv*. When *tvsl\_yield\_lv* is first called, *Complete\_Transversal* is an uninstantiated logic variable. Each time *tvsl\_yield\_lv* produces a result, *Complete\_Transversal* will have been unified with that result.

Section V-G discusses how *tvsl\_yield\_lv* maps to *tvsl\_prolog*.

### F. *tvsl\_prolog* (Listings in Appendix ??)

The final program, *tvsl\_prolog* (Listing ??), is straight Prolog. *tvsl\_prolog* and *tvsl\_yield\_lv* are the same program expressed in different languages. One can run *tvsl\_prolog* on, say, [SWI Prolog online](#) and get the result shown in Listing ??—although formatted somewhat differently.

## IV. CONTROL FUNCTIONS (LISTINGS IN APPENDIX ??)

This section discusses Prolog's control flow and explains how Pylog implements it. It also presents a number of Pylog control-flow functions.

### A. *Control flow in Prolog* (Listings in Appendix ??)

Prolog, or at least so-called "pure" Prolog, is a satisfiability theorem prover turned into a programming language. One supplies a Prolog execution engine with (a) a "query" or "goal" term along with (b) a database of terms and clauses and asks whether values for variables in the query/goal term can be found that are consistent with the database. The engine conducts a depth-first search looking for such values.

Once released as a programming language, programmers used Prolog in a wide variety of applications, not necessarily limited to establishing satisfiability.

An important feature of Prolog is that it distinguishes far more sharply than most programming languages between data flow and control flow.

- 1) By *control flow* we mean the mechanisms that determine the order in which program elements are executed or evaluated. This section discusses Pylog control flow.
- 2) By *dataflow* we mean the mechanisms that move data around within a program. Section V discusses how data flows through a Prolog program via logic variables and how Pylog implements logic variables.

The fundamental control flow control mechanisms in most programming languages involve (a) sequential execution, i.e., one statement or expression following another in the order in which they appear in the source code, (b) conditional execution, e.g., **if** and related statements or expressions, (c) repeated execution, e.g., **while** statements or similar constructs, and (d) the execution/evaluation of sub-portions of a program such as functions and procedures via method calls and returns.

Even declarative programming languages, such as Prolog, include explicit or implicit means to control the order of execution. That holds even when the language includes lazy

evaluation, in which an expression is evaluated only when its value is needed.

Whether or not the language designers intended this to happen, programmers can generally learn how the execution/evaluation engine of a programming language works and write code to take advantage of that knowledge. This is not meant as a criticism. It's a simple consequence of the fact that computers—at least traditional, single-core computers—do one thing at a time, and programmers can design their code to exploit that ordering.

Prolog, especially the basic Prolog this paper is considering, offers a straight-forward control-flow framework: lazy, backtracking, depth-first search. Listing ?? (See Bartak [2]) shows a simple Prolog interpreter written in Prolog. The code is so simple because unification and backtracking can be taken for granted!

The execution engine, here represented by the *solve* predicate, starts with a list containing the query/goal term, typically with one or more uninstantiated variables. It then looks up and unifies, if possible, that term with a compatible term in the database (line 3). If unification is successful, the possibly empty body of the clause is appended to the list of unexamined terms (line 4), and the engine continues to work its way through that list. Should the list ever become empty (line 1), *solve* terminates successfully. The typically newly instantiated variables in the query contain the information returned by the program's execution.

If unification with a term in the database (line 3) is not possible, the program is said to have *failed* (for the current execution path). The engine then backs up to the most recent point where it had made a choice. This typically occurs at line 3 where we are looking for a clause in the database with which to unify a term. If there are multiple such clauses, another one is selected. If that term leads to a dead end, *solve* tries another of the unifiable terms.

In short, terms either *succeed* in unifying with a database term,<sup>7</sup> or they *fail*, in which case the engine backtracks to the most recent choicepoint. This is standard depth-first search—as in *trvsl\_dfs\_first*. In addition, when the engine makes a selection at a choicepoint, it retains the ability to produce other possible selections—as in *tvsl\_yield*. The engine may be *lazy* in that it generates possible selections as needed.

Even when *solve* empties its list of terms, it retains the ability to backtrack and explore other paths. This capability enables Prolog to generate multiple answers to a query (but one at a time), just as *tvsl\_yield* is able to generate multiple transversals, but again, one at a time when requested.

Prolog often seems strange in that lazy backtracking search is the one and only mechanism Prolog (at least pure prolog) offers for controlling program flow. Although backtracking depth-first search itself is familiar to most programmers, lazy backtracking search may be less familiar. When writing Prolog code, one must get used to a world in which program flow is defined by lazy backtracking search.

<sup>7</sup>Operations such as arithmetic, may also fail and result in backtracking.

## B. Prolog control flow in Pylog (Listings in Appendix ??)

Prolog's lazy, backtracking, depth-first search is built on a mechanism that keeps track of unused choicepoint elements *even after a successful element has been found*. Let's compare the relevant lines of *tvsl\_dfs\_first* (Listing ??) and *tvsl\_yield* (Listing ??). We are interested in the **else** arms of these programs.

In both cases, the choicepoint elements are the members of *sets[0]*. (Recall that *sets* is a list of sets; *sets[0]* is the first set in that list. The choicepoint elements are the members of *sets[0]*.)

The first two lines of the two code segments are identical: define a **for**-loop over *sets[0]*; establish that the selected element is not already in the partial transversal.

The third line adds that element to the partial transversal and asks the transversal program (*tvsl\_dfs\_first* or *tvsl\_yield*) to continue looking for the rest of the transversal.

Here's where the two programs diverge.

- In *tvsl\_dfs\_first*, if a complete transversal is found, i.e., if something other than **None** is returned, that result is returned to the caller. The loop over the choicepoints terminates when the program exits the function via **return** on line 5.
- In *tvsl\_yield*, if a complete transversal is found, i.e., if **yield from** returns a result, that result is **yielded** back to the caller. But *tvsl\_yield* does *not* exit the loop over the choicepoints. The visible structure of the code suggests that perhaps the loop might somehow continue, i.e., that **yield** might not terminate the loop and exit the function the way **return** does. How can one return a value but allow for the possibility that the loop might resume? That's the magic of Python generators, the subject of the next section.

## C. A review of Python generators (Listings in Appendix ??)

This paper is not about Python generators. We assume readers are already familiar with them. Even so, because they are so central to Pylog, we offer a brief review.

Any Python function that contains **yield** or **yield from** is considered a generator. This is a black-and-white decision made by the Python compiler. Nothing is required to create a generator other than to include **yield** or **yield from** in the code.

So the question is: how do generators work operationally? Using a generator requires two steps.

- 1) Initialize the generator, essentially by calling it as a function. Initialization does *not* run the generator. Instead, the generator function returns a generator object. That generator object can be activated (or reactivated) as in the next step.
- 2) Activate (or reactivate) a generator object by calling *next* with the generator object as a parameter. When a generator is activated by *next*, it runs until it reaches a **yield** or **yield from** statement. Like **return**, a **yield** statement may optionally include a value to be returned

to the *next*-caller. Whether or not a value is sent back to the *next*-caller, a generator that encounters a **yield** stops running (much like a traditional function does when it encounters **return**).

Generators differ from traditional functions in that when a generator encounters **yield** it retains its state. On a subsequent *next* call, the generator resumes execution at the line after the **yield** statement.

In other words, unlike functions, which may be understood to be associated with a stack frame—and which may be understood to have their stack frame discarded when the function encounters **return**—generator frames are maintained independently of the stack of the program that executes the *next* call.

This allows generators to be (re-)activated repeatedly via multiple *next* calls.

Consider the simple example shown in Listing ?? . When executed, the result will be as shown in Listing ?? .

As *find\_number* runs through 1 .. 4 it **yields** them to the *next*-caller at the top level, which prints that they are not the search number. But note what happens when *find\_number* finds the search number. It executes **return** instead of **yield**. This produces a *StopIteration* exception—because as a generator, *find\_number* is expected to **yield**, not **return**. If the *next*-caller does not handle that exception, as in this example, the exception propagates to the top level, and the program terminates with an error code.

Python’s **for**-loop catches *StopIteration* exceptions and simply terminates. If we replaced the **while**-loop in Listing ?? with

```
for k in find_number(search_number):
    print(f {k} is not 5 )
```

the output would be identical except that instead of terminating with a *StopIteration* exception, we would terminate normally.

Notice also that the **for**-loop generates the generator object. The step that produces *find\_number\_object* (originally line 12) occurs when the **for**-loop begins execution.

**yield from** also catches *StopIteration* exceptions. Consider adding an intermediate function that uses **yield from** as in Listing ??.<sup>89</sup> The result is similar to the previous—but with no uncaught exceptions. See Listing ??.

Note that when *find\_number* fails in Listing ??, i.e., when *find\_number* does not perform a **yield**, the **yield from** line in *use\_yield\_from* does not perform a yield. Instead it goes on to its next line and prints the *find\_number failed* message. It then terminates without performing a **yield**, producing a *StopIteration* exception. The top-level **for**-loop catches that exception and terminates normally.

In short, because Python generators maintain state after performing a *yield*, they can be used to model Prolog backtracking.

<sup>8</sup>An intermediate function is required because **yield** and **yield from** may be used only within a function. We can’t just put **yield from** inside the top-level **for**-loop.

<sup>9</sup>This example was adapted from [this generator tutorial](#).

#### D. **yield** : *succeed* :: *return* : *fail* (Listings in Appendix ??)

Generators perform an additional service. Recall that Prolog predicates either *succeed* or *fail*. In particular when a Prolog predicate fails, it does not return a negative result—recall how *tvsl\_dfs\_first* returned **None** when it failed to complete a transversal. Instead, a failed predicate simply terminates the current execution path. The Prolog engine then backtracks to the most recent choicepoint.

Similarly, if a generator terminates, i.e., **returns**, before encountering a **yield**, it generates a *StopIteration* exception. The *next*-caller typically interprets that to indicate the equivalent of failure. In this way Prolog’s *succeed* and *fail* map onto generator **yield** and **return**. This makes it fairly straightforward to write generators that mimic Prolog predicates.

- A Pylog generator *succeeds* when it performs a **yield**.
- A Pylog generator *fails* when it **returns** without performing a **yield**.

Generators provide a second parallel construct. Multiple-clause Prolog predicates map onto a Pylog function with multiple **yields** in a single control path. The generic prolog structure as shown in Listing ?? can be implemented as shown in Listing ??.

Prolog’s **cut** (‘!’) (Listing ??) corresponds to a Python **if-else** structure (Listing ??). The two **yields** are in separate arms of an **if-else** construct.

The control-flow functions discussed in Section IV-E along with the *append* function discussed in Section V-E offer numerous examples.

Python’s generator system has many more features than those covered above. But these are the ones on which Pylog depends.

#### E. Control functions (Listings in Appendix ??)

Pylog offers the following control functions. (It’s striking the extent to which generators make implementation straightforward.)

- *fails* (Listing ??). A function that may be applied to a function. The resulting function succeeds if and only if the original fails.
- *forall* (Listing ??). Succeeds if all the generators in its argument list succeed.
- *forany* (Listing ??). Succeeds if any of the generators in its argument list succeed. On backtracking, tries them all.
- *trace* (Listing ??). May be included in a list of generators (as in *forall* and *forany*) to log progress. The second argument determines whether *trace* succeeds or fails. The third argument turns printing on or off. When included in a list of *forall* generators, *succeed* should be set to **True** so that it doesn’t prevent *forany* from succeeding. When included in a list of *forany* generators, *succeed* should be set to **False** so that *forany* won’t take *trace* as an extraneous success.
- *would\_succeed* (Listing ??). Like Prolog’s double negative,  $\backslash + \backslash +$ . *would\_succeed* is applied to a function. The resulting function succeeds/fails if and only if the original

function succeeds/fails. If the original function succeeds, this also succeeds but without binding any variables.

- *Bool\_Yield\_Wrapper*. A class whose instances are generators that can be used in **while**-loops. *Bool\_Yield\_Wrapper* instances may be created via a *bool\_yield\_wrapper* decorator. The decorator returns a function that instantiates *Bool\_Yield\_Wrapper* with the decorated function along with its desired arguments. The decorator is shown in Listing ??.

The example in Listing ?? uses *bool\_yield\_wrapper* twice, once as a decorator and once as a function that can be applied directly to other functions. The example also uses the *unify* function (see Section V-F below).

The output, as expected, is the first five squares.

Note that the **while**-loop on line 6 succeeds exactly once—because *unify* succeeds exactly once. The **while**-loop on line 10 succeeds 5 times.

An advantage of this approach is that it avoids the **for** loop. Notwithstanding our earlier discussion, **for**-loops don't feel like the right structure for backtracking.

A disadvantage is its wordiness. Extra lines of code (lines 4 and 9) to are needed to create the generator. One itches to get rid of them, but we were unable to do so.

Note that

```
while squares(5, Square).has_more():
```

does not work. The **while**-loop uses the entire expression as its condition, thereby creating a new generator each time around the loop.

Caching the generator has the difficulty that one may want the same generator, with the same arguments, in multiple places. In practice, we found ourselves using the **for** construct most of the time.

## V. LOGIC VARIABLES (LISTINGS IN APPENDIX ??)

Figure 1 shows Pylog's primary logic variable classes. This section discusses *PyValue*, *Var*, *Structure*, and the three types of sequences. (*Term* is an abstract class.)

### A. *PyValue* (Listings in Appendix ??)

A *PyValue* provides a bridge between logic variables and Python values. A *PyValue* may hold any immutable Python value, e.g., a number, a string, or a tuple. Tuples are allowed as *PyValue* values only if their components are also immutable.

### B. *Var* (Listings in Appendix ??)

A *Var* functions as a traditional logic variable: it supports unification.

Unification is surprisingly easy to implement. Each *Var* object includes a *next* field, which is initially **None**. When two *Vars* are unified, the *next* field of one is set to point to the other. (It makes no difference, which points to which.) A chain of linked *Vars* unify all the *Vars* in the chain.

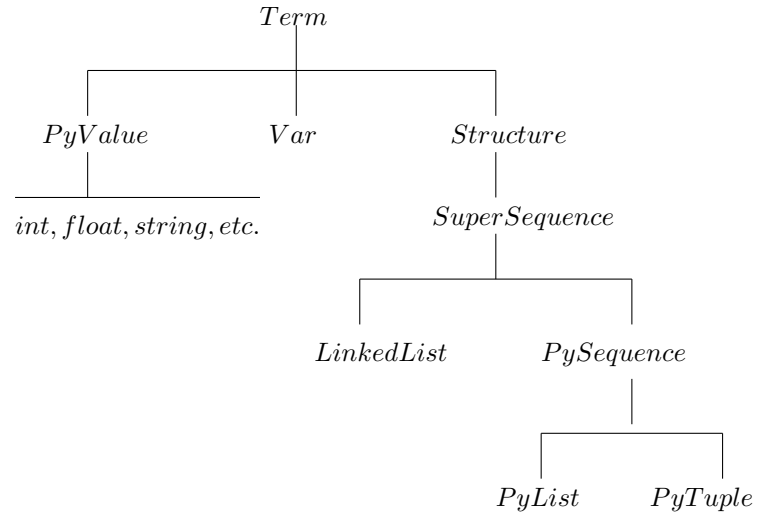


Fig. 1. This diagram shows a more complete list of Pylog classes.

Consider Listing ?. It's important not to be confused by **for**-loops. Even though the nested **for**-loops look like nested iteration, that's not the case. *There is no iteration!* In this example, the **for**-loops serve solely as choicepoints and scope definitions.

Since *unify* succeeds at most once, each **for**-loop offers only a single choice. There is never any backtracking. The only function of the **for**-loops is (a) to call the various *unify* operations and (b) to define the scope over which they hold.

The output (Listing ??) should make this clear.

Numbers with leading underscores indicate uninstantiated logic variables.

Line 1. All the logic variables are distinct. Each has its own identification number.

Line 2. *A* and *B* have been unified. They have the same identification number.

Line 3. *C* and *D* have also been unified. They have the same identification number, but different from that of *A* and *B*.

Line 4. All the logic variables have been unified—with a single identifier.

Line 5. All the logic variables have *abc* as their value.

Lines 6 - 9. Exit the unification scopes as defined by the **for**-loops and undo the respective unifications.

We can trace through the unifications diagrammatically. The first two unifications produce the following. (The arrows may be reversed.)

$$\begin{aligned} A &\rightarrow B \\ D &\rightarrow C \end{aligned} \tag{1}$$

The next unification is *A* with *C*. The first step in unification is to go to the end of the unification chains of the elements to



be unified. In this case,  $B$  (at the end of  $A$ 's unification chain) is unified with  $C$ . The result is either of the following.

$$\begin{array}{ccc} A & \rightarrow & B \\ & \downarrow & \\ D & \rightarrow & C \end{array} \qquad \begin{array}{ccc} A & \rightarrow & B \\ & \uparrow & \\ D & \rightarrow & C \end{array} \quad (2)$$

Finally, to unify  $E$  with  $D$ , we go to the end of  $D$ 's unification chain— $B$  or  $C$ .

$$\begin{array}{ccc} A & \rightarrow & B \\ & \downarrow & \\ D & \rightarrow & C \rightarrow E('abc') \end{array} \qquad \begin{array}{ccc} A & \rightarrow & B \rightarrow E('abc') \\ & \uparrow & \\ D & \rightarrow & C \end{array} \quad (3)$$

Different as they appear, these two structures are equivalent for unification purposes.

To determine a *Var*'s value, follow its unification chain. If the end is a *PyValue*, the *PyValue*'s value is the *Var*'s value. In (3), all *Vars* have value '*abc*'. If the end of a unification chain is an uninstantiated *Var* (as in (2) for all *Vars*), the *Var*'s in the tributary chains are mutually unified, but uninstantiated. When the end *Var* gets a value, it will be the value for all *Var*'s leading to it.

The following convenience methods make it possible to write the preceding code more concisely—but without the *print* statements. See Listing ??.

- *n\_Vars* takes an integer argument and generates that many *Var* objects.
- *unify\_pairs* takes a list of pairs (as tuples) and unifies the elements of each pair.

### C. Structure (Listings in Appendix ??)

The *Structure* class enables the construction of Prolog terms. A *Structure* object consists of a functor along with a tuple of values. The Zebra puzzle (Section VI) uses *Structures* to build *house* terms. *house* is the functor; the tuple contains the house attributes.

*house*(<nationality>, <cigarette>, <pet>, <drink>, <house color>)

*Structure* objects can be unified—but, as in Prolog, only if they have the same functor and the same number of tuple elements. To unify two *Structure* objects their corresponding tuple components must unify.

Let  $N$  and  $P$  be uninstantiated *Vars* and consider unifying the following objects.<sup>10</sup>

```
house(japanese, _, P, coffee, _)
house(N, _, zebra, coffee, _)
```

Unification would leave both *house* objects like this.

```
house(japanese, _, zebra, coffee, _)
```

Unification would have failed if the *house* objects had different *drink* attributes.

Prolog's unification functionality is central to how it solves such puzzles so easily. We discuss the *unify* function in Section V-F.

<sup>10</sup>The underscores represent don't-care elements.

### D. Lists (Listings in Appendix ??)

Pylog includes two *list* classes. *PySequence* objects mimic Python lists and tuples. They are fixed in size; they are immutable; and their components are (recursively) required to be immutable. The only difference between *PyList* and *PyTuple* objects is that the former are displayed with square brackets, the latter with parentheses.

More interestingly, Pylog also offers a *LinkedList* class. Its functionality is similar to Prolog lists. In particular, a *LinkedList* may have an uninstantiated tail, which is not possible with standard Python lists or tuples or with *PySequence* objects.

*LinkedLists* may be created in two ways.

- Pass the *LinkedList* class the desired head and tail, e.g.,  $Xs = \text{LinkedList}(Xs\_Head, Xs\_Tail)$ .
- Pass the *LinkedList* class a Python list. For example, *LinkedList*([]) is an empty *LinkedList*.

The next section (on *append*) illustrates the power of *LinkedLists*.

### E. append (Listings in Appendix ??)

The paradigmatic Prolog list function, and one that illustrates the power of logic variables, is *append/3*.

Pylog's *append* has Prolog functionality for both *LinkedLists* and *PySequences*. For example, running the code in Listing ?? produces the output in Listing ??.<sup>11</sup>

Pylog's *append* function for *LinkedLists* parallels Prolog's *append/3*. The Prolog code is in Listing ??; the Pylog code is in Listing ??.

Note that **yield from** appears twice. If after execution of the first **yield from** (line 3), *append* is called for another result, e.g., as a result of backtracking, it continues on to the second **yield from** (line 9). (As discussed in Section IV, this is standard behavior for Python generators.) The second part of the function calls itself recursively. Results are returned to the original caller from the first **yield from**—as in the Prolog version.

### F. Unification (Listings in Appendix ??)

To complete the discussion of logic variables, this section discusses the *unify* function—which, like so many Pylog functions, is surprisingly straightforward. (Listing ??.)

The *unify* function is called, *unify*(*Left*, *Right*), where *Left* and *Right* are the Pylog objects to be unified. (Argument order is immaterial.)

The first step (line 4) ensures that the arguments are Pylog objects. If either is an immutable Python element, such as a string or int, it is wrapped in a *PyValue*. This allows us to call, e.g., *unify*( $X$ , '*abc*') and *unify*('abc',  $X$ ).

There are four *unify* cases.

- 1) *Left* and *Right* are already the same. Since Pylog objects are immutable, neither can change, and there's nothing to do. Succeed quietly via **yield**.

<sup>11</sup>The output is the same whether we use *PySequences* or *LinkedLists*.

- 2) *Left* and *Right* are both *PyValues*, and exactly one of them has a value. Assign the uninstantiated *PyValue* the value of the instantiated one.

An important step is to set the assignment back to **None** after the **yield** statement. (line 18) This undoes the unification on backtracking.

- 3) *Left* and *Right* are both *Structures*, and they have the same functor. Unification consists of unifying the respective arguments.
- 4) Either *Left* or *Right* is a *Var*. Point the *Var* to the element at the end of the other element's unification chain. As line 1 shows, *unify* has a decorator. *euc* ensures that if either argument is a *Var* it is replaced by the element at the end of its unification chain. (*euc* stands for end of unification chain.) Again, unification must be undone on backtracking. (line 30)

#### G. Back to *tvsl\_yield\_lv* (Listings in Appendix ??)

We are now able to add more detail to our discussion of *tvsl\_yield\_lv* (Listing ??). We will step through the code line by line. We will see that *tvsl\_yield\_lv* is essentially a Pylog translation of *tvsl\_prolog* (Listing ??).

Line 2. *tvsl\_yield\_lv* has three parameters, as does *tvsl\_prolog*. (The other Python transversal programs had two.) The parameters of *tvsl\_yield\_lv* and *tvsl\_prolog* match up. In both cases. The third parameter is used to return the transversal to the caller.

Lines 3 and 4. These lines correspond to the second clause of *tvsl\_prolog*. (The first clause generates a log.) If we have reached the end of the sets, *Partial\_Transversal* is a complete transversal. Unify it with *Complete\_Tvsl*.

Lines 6-9. These lines correspond to the third clause of *tvsl\_prolog*.

Line 6 defines *Element* as a new *Var*.

Line 7 unifies *Element* with a member of *Sets[0]*. The Pylog *member* function is like the Prolog *member* function. On backtracking it unifies its first argument with successive members of its second argument. (This corresponds to line 10 of *tvsl\_prolog*.)

Line 8 ensures that the current value of *Element* is not already a member of *Partial\_Transversal*. (See the *fails* function in Section IV-E.) (This corresponds to line 11 of *tvsl\_prolog*.)

Line 9 calls *tvsl\_yield\_lv* recursively (via **yield from**). (This corresponds to lines 12 and 13 of *tvsl\_prolog*.)

## VI. THE ZEBRA PUZZLE (LISTINGS IN APPENDIX ??)

The Zebra Puzzle is a well known logic puzzle.

There are five houses in a row. Each has a unique color and is occupied by a family of unique nationality. Each family has a unique favorite smoke, a unique pet, and a unique favorite drink. Fourteen clues (Listing ??) provide additional constraints. *Who has a zebra and who drinks water?*

#### A. The clues and a Prolog solution (Listings in Appendix ??)

One can easily write Prolog programs to solve this and similar puzzles.

- Represent a house as a Prolog *house* term with the parameters corresponding to the indicated properties:

```
house(<nationality>, <cigarette brand>, <pet>, <
```

- Define the world as a list of five *house* terms, with all fields initially uninstantiated.
- Write the clues (Listing ??) as more-or-less direct translations of the English.

After the following adjustments, we can run this program online using SWI-Prolog.

- SWI-Prolog includes *member* and *nextto* predicates. SWI-Prolog's *nextto* means in the order given, as in clue 5.
- SWI-Prolog does not include a predicate for *next to* in the sense of clues 10, 11, and 14 in which the order is unspecified. But we can write our own, say, *next\_to*.

```
next_to(A, B, List) :- nextto(A, B, List).
next_to(A, B, List) :- nextto(B, A, List).
```

- Since none of the clues mentions either a zebra or water, we add the following.

```
% 15. (implicit).
member(house(_, _, zebra, _, _), Houses),
member(house(_, _, water, _, _), Houses).
```

When this program is run, we get an almost instantaneous answer—shown manually formatted in Listing ??.

*The Japanese have a zebra, and the Norwegians drink water.*

#### B. A Pylog solution (Listings in Appendix ??)

To write and run the Zebra problem in Pylog we built the following framework.

- We created a *House* class as a subclass of *Structure*. Users may select a house property as a pseudo-functor for displaying houses. We selected *nationality*.
- Each clue is expressed as a Pylog function. (See Listing ??.)
- The *Houses* list may be any form of *SuperSequence*.
- We added some simple constraint checking.

When run, the answer is the same as in the Prolog version. (See listing ??.)

Let's compare the underlying Prolog and Pylog mechanisms.

**Prolog.** It's trivial to write a Prolog interpreter in Prolog. See Listing ?? [2].

**Pylog.** We developed *three* Pylog approaches to rule interpretation.

- 1) *forall*. Use the *forall* construct as in Listing ??.  
*forall* succeeds if and only if all members of the list it is passed succeed. Each list element is protected within a **lambda** construct to prevent evaluation.
- 2) *run\_all\_rules*. We developed a Python function that accepts a list, e.g., of houses, reflecting the state of the world, along with a list of functions. It succeeds if and only if the functions all succeed. Listing ?? is a somewhat simplified version.
- 3) *Embed rule chaining in the rules*. For example, see Listing ??.  
Call *clue\_1* with a list of uninstantiated houses, and the problem runs itself.

The three approaches produce the same solution.

## VII. CONCLUSION (LISTINGS IN APPENDIX ??)

Embedding rule chaining in the clues suggests the template for Pylog as in Listing ??.

More generally, Pylog offers a way to integrate logic programming into Python.

- The magic of unification requires little more than linked chains.
- Prolog's control structures can be implemented as nested **for**-loops (for both choicepoints and scope setting), with **yield** and **yield from** gluing the pieces together.

## REFERENCES

- [1] Naveed Akhtar and Ajmal Mian. "Threat of adversarial attacks on deep learning in computer vision: A survey". In: *IEEE Access* 6 (2018), pp. 14410–14430.
- [2] Roman Bartak. *Meta-Interpreters*. in *Online Prolog Programming*. [http://kti.ms.mff.cuni.cz/~bartak/prolog/meta\\_interpret.html](http://kti.ms.mff.cuni.cz/~bartak/prolog/meta_interpret.html). 1998.
- [3] Shai Berger. *PYTHOLOGIC - PROLOG SYNTAX IN PYTHON (PYTHON RECIPE)*. <http://code.activestate.com/recipes/>. 303057-pythologic-prolog-syntax-in-python/. 2004.
- [4] Carl Friedrich Bolz. *A Prolog Interpreter in Python*. <http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.121.8625&rep=rep1&type=pdf>. 2007.
- [5] William E Byrd. *Relational programming in miniKanren: techniques, applications, and implementations*. 2010.
- [6] Bruno Kim Medeiros Cesar. *Prol: a minimal, inefficient Prolog interpreter in a few LOCs of Python*. <https://gist.github.com/brunokim/>. 2019.
- [7] Jan C Dageförde and Herbert Kuchen. "A compiler and virtual machine for constraint-logic object-oriented programming with Muli". In: *Journal of Computer Languages* 53 (2019), pp. 63–78.
- [8] Jan C Dageförde and Herbert Kuchen. "A constraint-logic object-oriented language". In: *Proceedings of the 33rd Annual ACM Symposium on Applied Computing*. 2018, pp. 1185–1194.
- [9] Christophe Delord. *PyLog*. <http://cdsoft.fr/pylog/index.html>. 2009.
- [10] Bruce Frederiksen. *Pyke*. <http://pyke.sourceforge.net/>. 2011.
- [11] Eugene C Freuder. "In pursuit of the holy grail". In: *Constraints* 2.1 (1997), pp. 57–61.
- [12] Jeremy Gibbons and Nicolas Wu. "Folding domain-specific languages: deep and shallow embeddings (functional pearl)". In: *Proceedings of the 19th ACM SIGPLAN international conference on Functional programming*. 2014, pp. 339–347.
- [13] Charles Antony Richard Hoare and He Jifeng. *Unifying theories of programming*. Vol. 14. Prentice Hall Englewood Cliffs, 1998.
- [14] Istasse Maxime. *Prology: Logic programming for Python3*. <https://github.com/mistasse/Prology>. 2016.
- [15] John McCarthy et al. "A Proposal for the Dartmouth Summer Research Project on Artificial Intelligence, August 31, 1955". In: *AI Magazine* 27.4 (Dec. 2006), p. 12. DOI: 10.1609/aimag.v27i4.1904. URL: <https://aaai.org/ojs/index.php/aimagazine/article/view/1904>.
- [16] Chris Meyers. *Prolog in Python*. <http://www.openbookproject.net/py4fun/prolog/prolog1.html>. 2015.
- [17] Nikola Miljkovic. *Python Prolog Interpreter*. <https://github.com/phonelines/Python-Prolog-Interpreter>. 2019.
- [18] Andrew Ng. *AI is the new electricity*. O'Reilly Media, 2018.
- [19] Ian Piumarta. *Lecture notes and slides from weeks 5-7 of a course on programming paradigms*. <http://www.ritsumei.ac.jp/~piumarta/pl/>. 2017.
- [20] Matthew Rocklin. *Kanren: Logic Programming in Python*. <https://github.com/logpy/logpy/>. 2019.
- [21] Stuart J Russell and Peter Norvig. *Artificial Intelligence-A Modern Approach, Third International Edition*. 2010.
- [22] Claudio Santini. *Pampy: The Pattern Matching for Python you always dreamed of*. <https://github.com/santinic/pampy>. 2018.
- [23] Ehud Y Shapiro. "The fifth generation project—a trip report". In: *Communications of the ACM* 26.9 (1983), pp. 637–641.
- [24] Ashish Shrivastava et al. "Learning from simulated and unsupervised images through adversarial training". In: *Proceedings of the IEEE conference on computer vision and pattern recognition*. 2017, pp. 2107–2116.
- [25] Herbert A Simon. "Models of man; social and rational." In: (1957).

- [26] Jeff Thompson. *Yield Prolog*. [http : / / yieldprolog . sourceforge.net/](http://yieldprolog.sourceforge.net/). 2017.
- [27] Markus Triska. “The Boolean Constraint Solver of SWI-Prolog: System Description”. In: *FLOPS: Functional and Logic Programming. 13th International Symposium*. Vol. 9613. LNCS. 2016, pp. 45–61.